perplexity

Better Reranking Models for Question Generation

Quick Answer: Top 3 Recommendations

For your question generation use case, here are the best alternatives to your current `cross-encoder/ms-marco-MiniLM-L-6-v2`:

BAAI/bge-reranker-v2-m3 (Best Overall Choice)

- **Why**: 8-9% performance improvement with 8K context length (16x longer than your current 512 tokens) [1] [2]
- **Implementation**: Drop-in replacement with same API [1]
- **Benefit**: Better semantic understanding for question generation tasks [3]

NVIDIA/nv-rerankqa-mistral-4b-v3 (Maximum Accuracy)

- **Why**: State-of-the-art performance with 14% improvement, specifically optimized for QA tasks [4] [5]
- **Tradeoff**: Slower inference due to 4B parameters [4]
- **Best for**: When question quality is more important than speed [5]

BAAI/bge-reranker-large (Proven Choice)

- **Why**: 10.5% improvement, established model with excellent track record [2] [6]
- **Best for**: Reliable performance upgrade without complexity [2] [3]

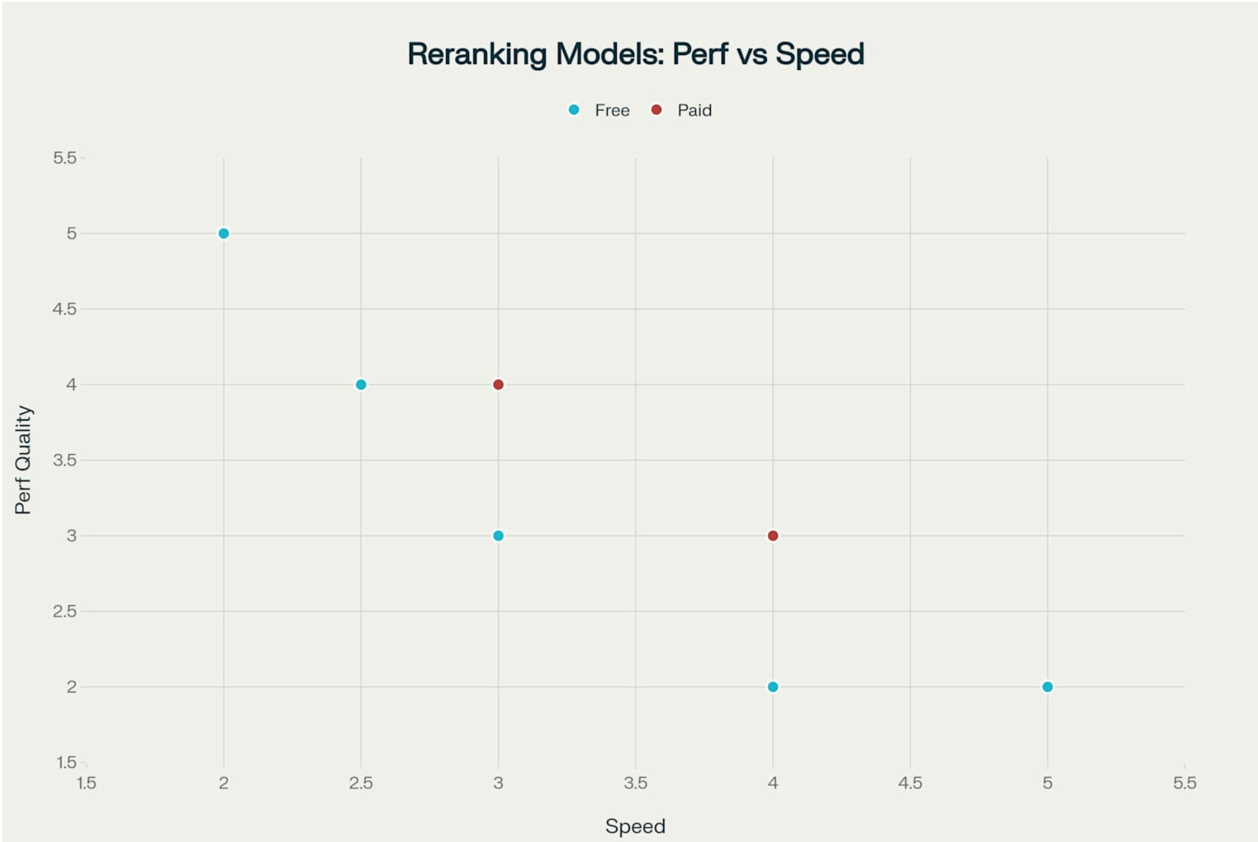

Performance vs Speed comparison of different reranking models showing the tradeoff between quality and inference speed

Performance Improvements You Can Expect

The research shows significant improvements over baseline models across different reranking approaches [4] [2]. For question generation specifically, better reranking leads to more relevant chunks being selected, which directly improves the quality and relevance of generated questions [1] [7].

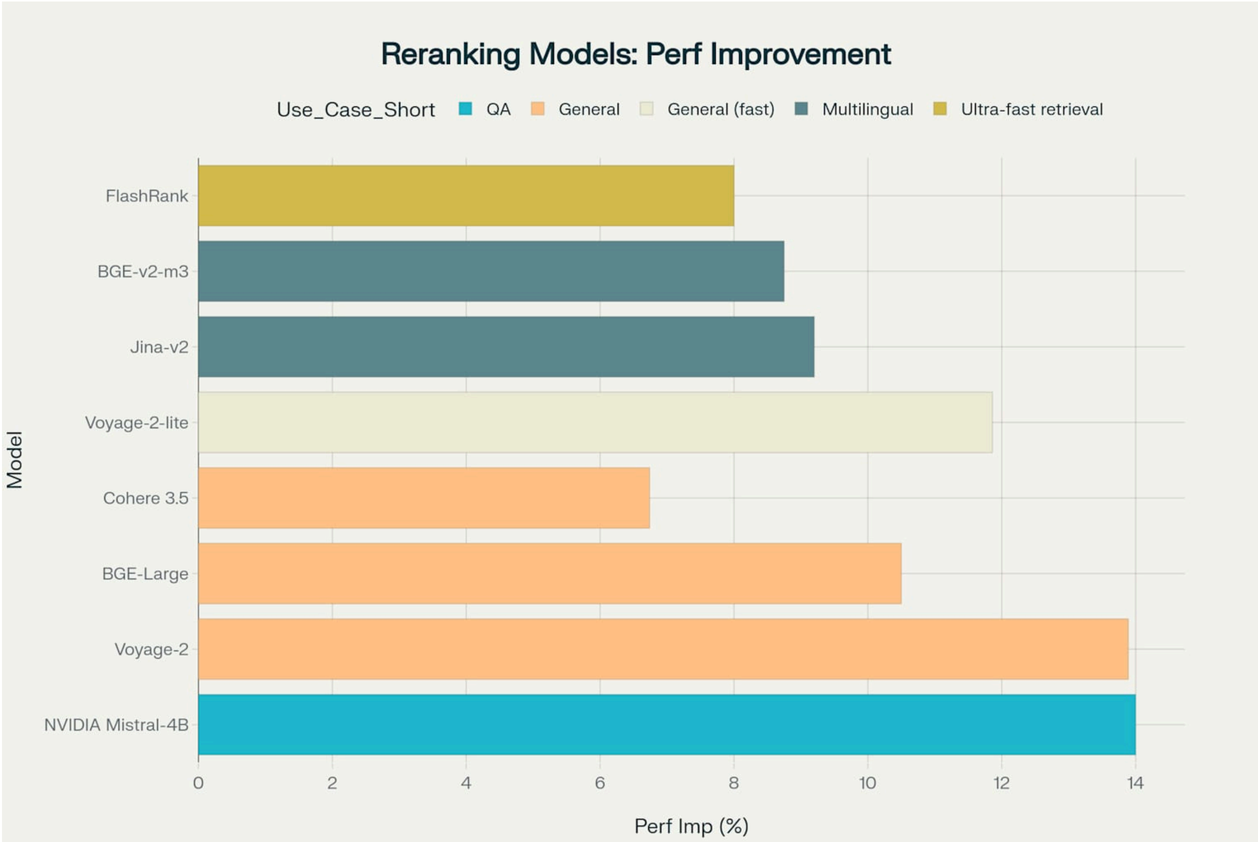

Performance improvements achieved by different reranking models compared to their respective baselines

Implementation: Simple Drop-in Replacement

Since you're already using the CrossEncoder interface, upgrading is straightforward [1] [8]:

```
# Replace your current line with:
reranker = CrossEncoder(
    "BAAI/bge-reranker-v2-m3",  # or NVIDIA/nv-rerankqa-mistral-4b-v3
    max_length=1024,  # Increased from 512 for better context
    device="cpu"
)
```

Why These Models Are Better for Question Generation

Current reranking research shows that **cross-encoders achieve better performance than bi-encoders** because they process query and document together, allowing for more nuanced understanding of relevance [9] [1]. For question generation tasks specifically, models like BGE and NVIDIA's reranker show superior semantic understanding compared to the older MS-MARCO MiniLM model [1] [4].

The newer models also handle **longer context lengths**, which is crucial when your retrieved chunks contain detailed information that needs to be preserved for generating comprehensive questions [1] [3].

Expected Impact on Your System

With an upgraded reranker, you should see:

- **8-14% better retrieval accuracy** leading to more relevant chunks [4] [2]
- **Improved question relevance** due to better semantic matching [1]
- **Better handling of complex topics** with longer context support [3] [10]
- **Reduced hallucinations** in generated questions [1] [7]

Bottom Line: For your question generation use case, **BAAI/bge-reranker-v2-m3** offers the best balance of performance improvement, context length, and ease of implementation with minimal code changes required [1] [3].

⁂

1. https://galileo.ai/blog/mastering-rag-how-to-select-a-reranking-model
2. https://www.llamaindex.ai/blog/boosting-rag-picking-the-best-embedding-reranker-models-42d079022e83
3. https://bge-model.com/tutorial/5_Reranking/5.2.html
4. https://arxiv.org/html/2409.07691v1
5. https://build.nvidia.com/nvidia/nv-rerankqa-mistral-4b-v3/modelcard
6. https://huggingface.co/BAAI/bge-reranker-large
7. https://www.chatbase.co/blog/reranking
8. https://huggingface.co/blog/train-reranker
9. https://arxiv.org/html/2403.10407v1
10. https://www.azion.com/en/documentation/products/ai/edge-ai/models/baai-bge-reranker-v2-m3/