

Chunk Count Versus Performance in Retrieval-Augmented Generation Systems

Increasing the number of chunks (top- k passages) that flow through a Retrieval-Augmented Generation (RAG) pipeline simultaneously improves recall while introducing latency, memory pressure, and noise that can eventually hurt answer quality^{[1] [2] [3]}. The optimal value of k is therefore a balance between evidence coverage and system efficiency, rather than “the larger the better.” This report synthesises recent empirical studies to explain how chunk count affects each stage of a modern RAG stack—from dense retrieval through cross-encoder reranking to sequence-to-sequence answer generation—and offers practical guidelines for setting k in production.


1 Background: What “Chunk Count” Means

Large corpora are first split into fixed-length passages (often 100-512 tokens) that are embedded and stored in a vector index^[4]. At query time the retriever returns the top k most similar chunks, which are optionally reranked and concatenated (or fused) into the language model’s context window for answer generation^{[5] [6]}. Here k is the chunk count whose impact we analyse.

2 Retrieval-Stage Effects

2.1 Recall Improves Logarithmically

Dense passage systems such as DPR and DensePhrases show sharp gains in top- k recall when k rises from 1 to 5, with ever-smaller improvements beyond 20^{[1] [4]}. For Natural Questions, DPR recall climbs from 44% at $k=1$ to 79% at $k=20$, while DensePhrases starts higher at 51% and plateaus earlier^[1].



Retrieval recall improves as k increases

2.2 Latency Grows Linearly

Although nearest-neighbour search is sub-millisecond, ranking or merging thousands of candidates can add noticeable overhead, especially when hybrid queries call both dense and sparse indices^{[7] [3]}. Empirical cost models report retrieval latency roughly proportional to k until GPU memory bandwidth becomes a bottleneck^[3].

2.3 Risk of Semantic Drift

Adding many marginally relevant passages raises the probability that the reranker or reader focuses on distractors, a phenomenon called “detrimental context”^[8]. Studies show answer accuracy can **drop by up to 10%** when irrelevant documents dominate the top- k list in large collections^{[9] [8]}.

3 Reranking-Stage Effects

3.1 Cross-Encoder Capacity

Cross-encoder rerankers (e.g., MiniLM or BGE) score every query–chunk pair, so their compute scales with $k \times \text{sequence length}$ ^[5]. MiniLM can rerank 1 800 chunks / s on a single GPU at 512 tokens, whereas larger multilingual models slow to 400-600 chunks / s^[3].

3.2 Precision Saturation

Beyond 30–50 candidates, rerankers rarely reorder documents enough to change the reader’s final answer^{[10] [11]}. Therefore many pipelines rerank only the first 50–100 hits even when the retriever produced 1 000 candidates^[11].


3.3 Noise-Aware Rerankers

Recent methods (e.g., KG-FiD and MGFiD) jointly learn to drop low-value chunks early, recovering FiD-level accuracy while processing just 20 passages—cutting computation by 40–60%^{[5] [2]}.

4 Reader-Stage Effects

4.1 Answer Accuracy versus k

Fusion-in-Decoder (FiD) accuracy rises steeply between $k=5$ and $k=25$, then plateaus^[5]. On Natural Questions, EM jumps from 38% at $k=5$ to 45% at $k=25$, with only a 0.4% gain when k doubles to 50^[5].



Impact of increasing top-k passages on answer accuracy

4.2 Context Window Saturation and Token Limits

When k grows, concatenated passages rapidly exhaust the model's context window. Long-context LLMs can ingest 64 k tokens, yet experiments show many models degrade after 32 k because attention sparsity becomes a bottleneck or the network fails to prioritise relevant spans^{[12] [13]}.

4.3 Memory-Bandwidth Bottleneck

Decoder cross-attention time dominates FiD inference once passages exceed 20–40. Profiling reveals that **over 60%** of end-to-end latency occurs in the decoder despite encoder FLOPs being higher, because cross-attention is memory-bandwidth bound^[3].

4.4 Noise Amplification

Adding more passages that partially mention the answer can mislead the generator. Filtering out noisy chunks after the 12th encoder layer lets KG-FiD match full FiD accuracy while evaluating only 20 passages, demonstrating that intelligent pruning outperforms blind increases in k ^[5].

5 Resource Trade-offs

Metric	Retrieval	Reranking	Reader
FLOPs growth	$O(k)$	$O(k)$	$O(k \cdot L^2)$
GPU memory	Vector cache scales with $k \times \text{dim}$ ^[3]	negligible	Cross-attention key-value pairs scale with $k \times \text{tokens}$ ^[3]
Latency tipping-point	~100–200 passages	~50–100 passages	~20–40 passages ^[3]
Typical production k	20–40	50–100 (but prune after score)	≤ 30 ^[3]

6 Practical Guidelines

- **Start with $k=20$.** This captures most recall gains while keeping cross-attention manageable^{[1] [5]}.
- **Measure marginal EM gains.** If increasing k from 20 to 40 boosts EM by <1%, keep the smaller value^{[5] [8]}.
- **Use staged pruning.** Retrieve 1 000, rerank top 100, read top 30; this yields near-oracle coverage with acceptable cost^{[11] [5]}.
- **Adopt noise-aware models.** Techniques like END or passage classifiers drop distracting chunks, letting you raise recall without length penalties^{[2] [14]}.
- **Profile latency under load.** Decoder time dominates at high k ; FiD-Light or FiDO compress encodings to triple throughput with minor accuracy loss^[3].
- **Monitor context length limits.** Even 200 k-token models degrade beyond 64 k, so rely on smarter retrieval, not blindly larger k ^[12].

7 Conclusion

Increasing the number of retrieved chunks initially improves both retrieval recall and answer quality, but returns diminish and can reverse once noise and computational overhead outweigh the benefits^{[1] [8] [3]}. State-of-the-art systems therefore combine moderately sized k (≈ 20 – 40) with learned rerankers and context denoising to achieve a sweet spot of accuracy, cost, and latency. Treat k as a tunable hyper-parameter—scaled dynamically with query complexity and hardware budget—rather than a fixed target to maximise.

✱

1. <https://stackoverflow.com/questions/63850906/top-k-precision>
2. <https://arxiv.org/pdf/2308.04028.pdf>
3. <https://milvus.io/ai-quick-reference/how-can-we-measure-the-accuracy-of-the-retrieval-component-in-a-rag-system-for-example-using-metrics-like-precisionk-and-recallk-on-the-documents-retrieved>
4. <https://trec.nist.gov/pubs/trec29/papers/UoB.DL.pdf>
5. <https://milvus.io/ai-quick-reference/what-is-the-optimal-chunk-size-for-rag-applications>
6. <https://qdrant.tech/documentation/beginner-tutorials/retrieval-quality/>
7. <https://aclanthology.org/2021.findings-acl.29.pdf>
8. <https://milvus.io/ai-quick-reference/what-is-an-acceptable-latency-for-a-rag-system-in-an-interactive-setting-eg-a-chatbot-and-how-do-we-ensure-both-retrieval-and-generation-phases-meet-this-target>
9. <https://vectorize.io/top-5-metrics-to-master-for-effective-rag-assessments/>
10. <https://arxiv.org/html/2404.07220v2>
11. <https://www.datacamp.com/tutorial/how-to-improve-rag-performance-5-key-techniques-with-examples>
12. <https://www.linkedin.com/pulse/enhancing-llm-accuracy-retrieval-augmented-generation-mark-hinkle-papxe>
13. <https://arxiv.org/html/2407.02485v1>
14. <https://aclanthology.org/2022.emnlp-main.249.pdf>